

“Concurrency in Go” Notes

BY ARNAV KUMAR

Web: <https://arnavcs.github.io>

Concurrency in Go is a publication by O'Reilly Media Inc. written by Katherine Cox-Buday. This is a collection of notes that I make about the text as I read it, and is not a summary or recreation of the text, but rather a reference for anyone who has already read the text. As such, please read the text to gain a better understanding of the contents.

1 Basic Concurrency Ideas

Color Scheme Key
Software / Design Pattern
Definition
Note

Amdahl's Law

Amdahl's Law models the improved performance of a fixed task when the resources are improved. In parallel computing, it is used to predict the speedup of using multiple processors. The relation is given as follows:

$$S_{\text{latency}}(s) = \frac{1}{(1-p) + \frac{p}{s}}$$

S_{latency}	the theoretical speedup of the whole program
s	the speedup of the part of the task from improved resources
p	the proportion of the execution time that benefits from the improved resources

Race Conditions

A race condition is when two or more operations must execute in the correct order, but the program leaves the order of execution unspecified.

Data Race

A data race is a race condition in which two concurrent operations attempt to read the same data at an unspecified time (namely one that could potentially conflict). In the following example, the program is not given a specified evaluation order, so the code that follows may execute before, during, or even after the goroutine. As such, the output is indeterminate.

```
var data int
go func() { data++ }()
fmt.Printf("%v\n", data)
```

Atomicity

An atomic operation is indivisible or uninteruptable in the context in which it is operating. For example, the statement `i++` consists of 3 atomic operations: retrieving, incrementing, and storing the value of `i`.

Critical Selection	<p>A critical selection is a section of code that requires exclusive access to a shared resource. In the following code, the <code>fmt.Printf()</code> and the goroutine are both critical selections.</p> <pre>var data int go func() { data++ }() fmt.Printf("%v\n", data)</pre>								
Memory Access Synchronization	To solve the problem of multiple critical selections, only enable one critical selection to access the same shared resource at a time. This can be achieved, for example, with a mutex.								
Deadlock	A deadlock is a state in which all concurrent processes are waiting on each other. A deadlock can be identified by the Coffman Conditions.								
Coffman Conditions	<p>There are 4 Coffman Conditions that detect, prevent, and correct deadlocks. The conditions are as follows:</p> <table> <tr> <td>Mutual Exclusion</td><td>A concurrent process must hold exclusive rights to a resource at any one time.</td></tr> <tr> <td>Wait For Condition</td><td>A concurrent process must hold a resource and be waiting for another resource.</td></tr> <tr> <td>No Preemption</td><td>A resource held by a concurrent process can only be released by that process.</td></tr> <tr> <td>Circular Wait</td><td>A process must be waiting on a chain of processes which is circular (meaning that the process is directly or indirectly waiting on itself to give a result).</td></tr> </table>	Mutual Exclusion	A concurrent process must hold exclusive rights to a resource at any one time.	Wait For Condition	A concurrent process must hold a resource and be waiting for another resource.	No Preemption	A resource held by a concurrent process can only be released by that process.	Circular Wait	A process must be waiting on a chain of processes which is circular (meaning that the process is directly or indirectly waiting on itself to give a result).
Mutual Exclusion	A concurrent process must hold exclusive rights to a resource at any one time.								
Wait For Condition	A concurrent process must hold a resource and be waiting for another resource.								
No Preemption	A resource held by a concurrent process can only be released by that process.								
Circular Wait	A process must be waiting on a chain of processes which is circular (meaning that the process is directly or indirectly waiting on itself to give a result).								
Livelock	A livelock is when the current concurrent processes are performing operations, but these operations do not terminate or move the program closer to termination.								
Starvation	Starvation is a superset of a livelock or deadlock where, more generally, a concurrent process does not receive access to the resources it needs. A common example is having a “greedy worker” hold on to access to the resource, while a “polite worker” does not, and thus has less access to the resource: it is starved.								
“Finding a Balance”	What should the range of a memory lock be? Should it be broad and cover multiple critical selections, or should each critical selection get its own lock? It is important to strike a balance in answering this question because memory access synchronization is expensive, but you also want to avoid writing greedy processes to mitigate starvation.								
OS Threads	OS threads are a primitive at the OS context that can be used to run processes concurrently. The operating system is responsible for creating and managing the threads. The threads all have access to a shared resource space.								
Green Threads	Green threads are threads that are managed by a program’s runtime.								
Coroutines	Concurrent subroutines that are non-preemptive (meaning that they can’t be interrupted) are called coroutines. They feature multiple points to suspend or reenter computation.								

M:N Scheduler	A M:N scheduler is the mechanism that Golang uses to host goroutines and it consists of mapping M green threads onto N OS threads.
Fork-Join Model	The model that Golang follows for concurrency, a fork-join model is one in which a child branch can fork off from parent to be run concurrently. After the termination of the child branch, it is joined back to the parent branch at a join point.
Thread Pools	Thread pools are a software design pattern that maintains a collection of threads to map incoming tasks to threads for concurrent execution.
Concurrency vs. Parallelism	<p>Parallelism is a property of a machine to be able to run two tasks simultaneously in the considered context. On the other hand, concurrency refers to when two processes have a lifespan that overlaps. In this sense, you could have a concurrent program running on a single thread where multiple threads are simulated. It is also possible that the concurrent processes run in parallel.</p> <p>Concurrency is a property of the code, and parallelism is a property of the execution of the code.</p>
Process	A process is a portion of code that requires input to run and produces an output that is consumed by another process. The input and output of a process is called communication between processes.
Communicating Sequential Processes (CSP)	CSP is the name of a paper, programming language, and the idea of a describing programs as processes which are sequential and communicate. Used in the paper describing CSP, the CSP language supported the use of ! and ? to send input into and read output from a process respectively. In addition, it supported guarded commands. This is the style of concurrency programming that Golang's channels are based on.
Guarded Command	<p>When a statement should not be executed if another statement was false or a command exited, it is a guarded command. The CSP example below denotes a process a, from which a character c is continually read (while there is something to be read), and then inputted into the process b.</p> <pre>*[c:character; a?c -> b!c]</pre>
Process Calculus	Process calculus is a mathematical way to model concurrent systems and analyze their properties.
Should I use CSP style or OS threads?	<p>The CSP style has certain advantages that it comes with, and more generally, the Golang developing team suggest to use the CSP style over primitives like <code>sync.Mutex</code>, but there are certain guidelines outlined that help determine when you should use channels or OS thread primitives. Follow the first applicable statement.</p> <ol style="list-style-type: none"> 1. If your code is performance critical, use primitives 2. If you are trying to transfer ownership of data, use channels 3. If you are trying to guard the internal structure of a struct, use primitives 4. If you are coordinating multiple pieces of logic, use channels 5. Use primitives

Mutex	Mutex stands for “mutual exclusion” and enables a way to express exclusive access to a shared resource. A mutex is often used for critical selections.
Object Pool	This pattern is a way to create a fixed number of objects for use, and is especially useful for objects that are computationally expensive or objects that will take a lot of memory.
Channels	<p>The channel pattern is a way to pass information. If there is nothing to be read from a channel, reading from it blocks execution; waiting for a value to be added to the channel. Additionally, channels can be closed (to stop writing to the channel), in which case reading from the channel further does not return a meaningful value, but will indicate that the channel is closed. Channels can also have buffers to store values to be read later.</p> <p>As a pattern, to write robust code, separate the ownership of the channel so that the channel utilizers only have read access to the channel, and the channel owner has the following responsibilities:</p> <ol style="list-style-type: none"> 1. Instantiate the channel 2. Perform writes or pass write ownership to another goroutine 3. Close the channel 4. Expose a reader channel for the channel utilizers

2 Golang Features and Building Blocks

Color Scheme Key
Type
Function / Keyword

func	<p>This keyword can be used to create named functions, closures, or anonymous functions. A named function example is show below.</p> <pre>func helloWorld(numTimes int) { for ; numTimes > 0; numTimes-- { fmt.Printf("Hello World!\n") } }</pre> <p>An anonymous version of the same function is also shown below.</p> <pre>var f := func(numTimes int) { for ; numTimes > 0; numTimes-- { fmt.Printf("Hello World!\n") } }</pre>
Loops	
range	

go

Creates a *goroutine* that runs the function, method, or closure concurrently by multiplexing onto OS threads. Each goroutine is a special class of coroutine where you do not have to manually describe the suspension and resuming of the routine. At runtime, Golang automatically suspends goroutines when they are blocked and resumes them when they are unblocked. Goroutines use the fork-join model for concurrency and during runtime, a M:N scheduler is used. See the following example using goroutines modified from the textbook that uses closures to print "go", "rust", and "c" concurrently in an unspecified order.

```
var wg sync.WaitGroup
for _, lang := range []string{"go", "rust", "c"} {
    wg.Add(1)
    go func(l string) {
        defer wg.Done()
        fmt.Println(l)
    }(lang)
}
wg.Wait()
```

defer

Defers the execution of the statement to the end of the function. In the following example, the mutex isn't unlocked until after the value of `data` increments.

```
var data int = 0
var mu sync.Mutex
func inc() {
    mu.Lock()
    defer mu.Unlock()
    data++
}()
```

struct

A struct in Golang can be named or anonymous. To create a named struct, follow the following syntax.

```
type Fruit struct {
    name string
}
var apple Fruit = Fruit{"Apple"}
```

And to create an anonymous struct, the following syntax holds.

```
apple = struct {name string} {name: "Apple"}
```

interface

Interfaces in Golang can be declared as follows.

```
type Plant interface {
    getHeight() float
    getSpecies() string
}
```

Additionally, the existence of the empty `interface` in Go is special, because all types satisfy the empty interface, meaning it can hold any value. It is `interface{}`.

sync.WaitGroup

A waitgroup stops the execution of certain code past a point until all processes being waited on are completed. It supports methods such as `.Add(int)`, `.Wait()`, and `.Done()`. You should use a `WaitGroup` when you do not care about the results of the concurrent operations or have another mean to collect the results. If that is not the case, use a `select` statement with channels. See the example below which prints the number from 0 to n-1 in some unspecified order.

```
const n = 3
var wg sync.WaitGroup
wg.Add(n)
for i := 0; i < n; i++ {
    go func(wg *sync.WaitGroup, i int) {
        defer wg.Done()
        fmt.Println(i)
    }(&wg, i)
}
wg.Wait()
```

sync.Mutex

A mutex type that supports the `.Lock()`, `.TryLock()`, and `.Unlock()` methods. These methods declare exclusive access to the shared resource that the mutex represents. By convention, a mutex unlock statement is in a `defer` statement to avoid **panicing** meaning that the mutex is not unlocked.

sync.RWMutex

This form of mutex requires the specification of the type of access desired. An arbitrary number of readers are allowed to read the same resource granted that there are no writers. In exchange for the greater control over the memory (and potentially less opportunity for starvation), it gives lower performance than `sync.Mutex` for a small number of readers. When the number of readers is high, though, its performance is noticeable. The supported methods are those from `sync.Mutex`, and the additional `.RLock()`, `.TryRLock()`, `.RUnlock()`, and `.RLocker()`.

`sync.Cond`

A `sync.Cond` is a “rendevous point” for goroutines waiting for an event (an signal between two or more goroutines that carries no information). The instantiation of a `Cond` is done with `sync.NewCond` which takes a `sync.Locker` interface (accessible with `.L`). Additionally, the methods `.Broadcast()`, `.Signal()`, and `.Wait()` are available to be used. Consider the following function from the textbook that “subscribes” a function to a `Cond`, running the function once when the `Cond` first broadcasts.

```
subscribe := func(c *sync.Cond, f func()) {
    var goroutineRunning sync.Waitgroup
    goroutineRunning.Add(1)
    go func() {
        goroutineRunning.Done()
        c.L.Lock()
        defer c.L.Unlock()
        c.Wait()
        f()
    }()
    goroutineRunning.Wait()
}
```

`sync.Once`

A variable, `once`, of type `sync.Once` will support the `.Do(func())` method which will only execute the passed function once regardless of what if a different function is passed.

`sync.Pool`

A pool object is an implemenation of an object pool. It can be instantiated by specifying the `New` field which is a thread safe member variable function that creates a new object in the pool. the `Pool` also supports the methods `.Get()`, and `.Put(object)`. Make no assumptions about the state of the instance you get back from `.Get()`, but objects in the `Pool` should be roughly uniform in makeup.

Channels

Channels that are read-write, read-only, and write-only that carry values of type `T` have types `chan T`, `<-chan T`, and `chan<- T` respectively. A channel can be closed if it is writable, and is done so with `close()`. To read all the values in the channel until it is closed, use `range`. Additionally, buffer size of the channel can be specified during initiation, and the default buffer size is 0. Reading from a channel instantiated with a buffer of capacity 4 can look as follows.

```
channelOwner := func() <-chan int {
    intStream := make(chan int, 4)
    go func() {
        defer close(intStream)
        for i := 0; i < 10; i++ { intStream <- i }
    }()
    return intStream
}

readIntStream := channelOwner()
for element := range readIntStream {
    fmt.Println(element)
}
```

`make()`

`new()`

Type Assertions

`select`

The `select` statement is able to bind channels together. Namely, all `case` statements are simultaneously checked to see if they are ready (for reading this is a populated or closed channel, and for writing this is a channel not at capacity). If there is no `default`, then the execution is blocked until one of the channels is ready. One of the cases is then chosen at random, and the associated statements run. If there is a `default`, then the execution isn't blocked. This can be used to complete other tasks while waiting for a result. See the example below modified from the textbook which demonstrates using a `select` statement to complete work while waiting.

```
done := make(chan interface{})
go func() {
    defer close(done)
    time.Sleep(5 * time.Second)
}()
loop:
for {
    select {
    case <-done:
        break loop
    default:
        workCounter++
        time.Sleep(time.Second)
    }
}
fmt.Println(workCounter)
```

`runtime.`
`GOMAXPROCS()`

Takes an integer parameter that specifies the number of OS threads that will host “work queues”.

3 Concurrency Patterns in Golang